

Held-out gap by target level and search strategy

$|\rho_{\text{hat_task}} - \rho|$ on a held-out instantiation of each strategy's selected configuration

Declared target surrogate tier_mid (c=0.55) - 256 tasks per instantiation - 10 iterations per run - no language model executed



Mean gap across the four levels

● Feedback (BeTaL-style)

search 5.23% held-out 2.87%

■ RS+PPR baseline

search 20.90% held-out 13.61%

▲ Best-of-N baseline

search 23.72% held-out 11.46%

Lower is better. Values are exact and reproducible from per_level_metrics.csv; no confidence intervals were run.